03 - Exploratory Data Analysis of Point Cloud Features

May 5, 2021

1 Exploratory Data Analysis of Point Cloud Features

In this notebook, an exploratory data analysis for the features that are calculated from the 3D point cloud (see previous exercise notebook) is performed.

The learning objectives of this notebook are:

- Explore the features of the 3D point cloud and calculated features
- Hand-on experience with plots
- Get a brief introduction by example to matplotlib

Please note that there are many different ways to generate one and the same plot using matplotlib. The ones shown in this notebook uses the object-oriented interface and might not be the shortest way to achieve the result. There are often shortcuts that can be used.

For a more in-depth coverage of matplotlib, we refer to the literature in ISIS. This notebook is not intended to be a comprehensive and complete documentation of matplotlib.

2 Point properties

We start exploring the original 3D point cloud that was used for feature extraction. Especially the point properties recorded by the laser scanner sensor are of interest.

Load the original 3D point cloud file to explore the stored information therein.

```
[1]: from pathlib import Path
import os

data_dir = str(Path.home()) + r'/coursematerial/GIS/ISPRS'

filepath = os.path.join(data_dir, r'Vaihingen3DTraining.las')

print(filepath)
```

/home/jovyan/coursematerial/GIS/ISPRS/Vaihingen3DTraining.las

```
[2]: import laspy import numpy as np
```

```
file = laspy.file.File(filepath, mode='r')
```

2.1 Classification

The 3D point cloud contains a classification of the individual points according to the following classes:

Power line, Low vegetation, Impervious surfaces, Car, Fence/Hedge, Roof, Facade, Shrub, Tree We assign the NumPy array of the LAS file to the variable classes.

```
[3]: classes = file.classification
```

And define a list of class names.

```
[4]: class_names = ['Powerline', 'Low vegetation', 'Impervious surfaces', 'Car', 

→ 'Fence/Hedge', 'Roof', 'Facade', 'Shrub', 'Tree']
```

With the function **unique()**, we can get the unique values of the array as well as the number of times this unique value is present.

```
[5]: u, c = np.unique(classes, return_counts=True)
    print(u)
    print(c)
```

```
[0 1 2 3 4 5 6 7 8]
[ 546 180850 193723 4614 12070 152045 27250 47605 135173]
```

For a nicer output, we use pandas DataFrames, which we construct from the classes array and the list of class names as the index. (For a more detailed explanation of pandas, please refer to the respective notebook. At this point, it is not really necessary to understand pandas well.)

```
[6]:
                            labels
     Powerline
                               546
     Low vegetation
                            180850
     Impervious surfaces
                           193723
     Car
                              4614
     Fence/Hedge
                             12070
     Roof
                            152045
     Facade
                             27250
     Shrub
                             47605
```

Tree 135173

As we can see from the output, there are only few points of the classes power line, car, fences, facades, and maybe shrub.

For categorical data, a bar plots can be used to show the classes.

After using the magic command (%matplotlib inline) and importing the library, a figure of a certain size is constructed with the **figure()** function. A figure is top-level container for all the plot elements. And you can set figure properties that are valid for the whole figure and not just for parts thereof like the size of the figure or the resolution in DPI (dots per inch). With the method **suptitle()**, we give the figure a title.

Using the figure object, an axes is added to the figure with the method add_axes(), providing the size of a rectangle that the axes should take within the figure. An axes is again a container that holds a specific type of plot. The rectangle parameter is a quadruple with the elements left, bottom, width, and height. Left and bottom are the position of the axes within the figure, and width and height the size. The values of all four dimensions are given relative to the size of the figure with values between 0.0 and 1.0. By adding several axes objects to a figure, you can layout them with the rectangle parameter. But often it is sufficient to use some convenient methods for constructing a figure with several axes. (Try changing the parameter values for the add_axes() method to, e.g., 0.5 or 0.3 and see what happens.)

The axes object is then used to construct a bar plot with the method **bar()** taking the x-coordinates values as first parameter, the heights as second parameter, and then tick labels with the named **tick_label** parameter. For our plot, the unique values (u) are the x-coordinates, the count values (c) the y-coordinates, and the class names the tick labels.

If we do not change the rotation of the tick labels, then the class names are written on top of each other as there is not enough horizontal space under the plot. We therefore change the rotation of all tick labels on the x-axis to a vertical rotation. We therefore loop through all tick labels of the x-axis with the method **xticklabels()**, and set the rotation of each tick label to vertical with the method **set_rotation()**. Instead of the string 'vertical', you can also provide a rotation angle in degrees.

Finally, the labels of the x-axis and the y-axis are set with the methods **set_xlabel()** and **set_ylabel()**.

```
[7]: %matplotlib inline

import matplotlib.pyplot as plt

# construct a figure of a certain size
fig = plt.figure(figsize=(8, 5))

# set title of figure
fig.suptitle('Points per Class', fontsize=14, fontweight='semibold')

# add an axes that takes almost all the space of the figure,
# but leaves some space at the top for the title
```

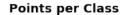
```
ax = fig.add_axes((0, 0, 1, 0.90))

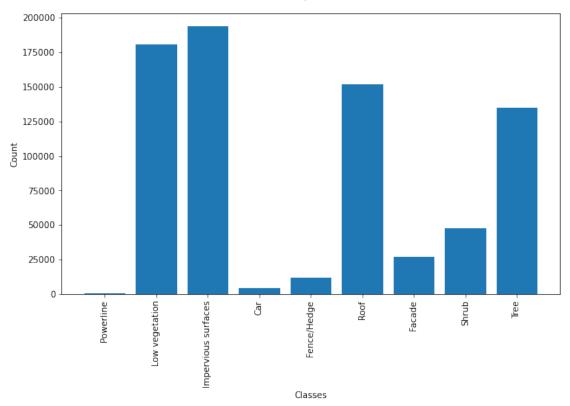
# draw a bar plot with u as the difference x values,
# c as the height, and class names for the labels
ax.bar(u, c, tick_label=class_names)

# rotate all labels to vertical (rotate by 90 degree)
for tick in ax.get_xticklabels():
    tick.set_rotation('vertical')

# set labels of axes
ax.set_xlabel('Classes')
ax.set_ylabel('Count')
```

[7]: Text(0, 0.5, 'Count')

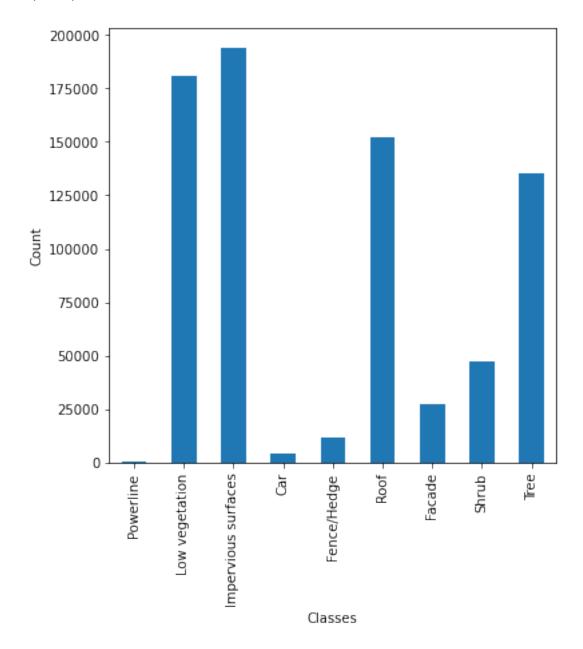




The pandas DataFrame class also provides a lot of convenient plotting methods. Pandas uses internally the matplotlib, so that the **bar()** method returns the axes object, which can be used to change some properties of the plot.

```
[8]: ax = df.plot.bar(figsize=(6,6), legend=False)
ax.set_xlabel('Classes')
ax.set_ylabel('Count')
```

[8]: Text(0, 0.5, 'Count')



From the plot, the absolute number of points per class, but also the proportions between the individual classes, become readily apparent. From a machine learning perspective, the underrepresented classes will probably be difficult to learn. But for low vegetation, impervious surfaces, roof, and tree, there should be enough points.

The LAS file also contains information about **intensity**, **number of returns** (**num_returns**), and **return number** (**return_num**), which we explore in the following.

Intensity is the returned intensity of the laser beam. The laser emits a laser signal with a certain intensity. This laser signal hits an object and is partially absorbed and partially reflected, where the ratio depends on the material of the surface. The portion of the laser beam that is reflected and reaches the laser scanner again is recorded as the intensity value.

Some objects reflect the laser signal only partially, so some portion is reflected and another portion is going through the object. This typically happens when the laser hits higher vegetation like trees, where the signal is then reflected at different height levels at the branches until it reaches the ground. Then we have several returns. The **number of returns** (**num_returns**) is the number of returns the laser scanner recorded for this measured point.

If there are several returns per laser beam, then these returns are converted into different point coordinates. This is because the laser beam traveled further towards the ground than the reflected signal and reached an object at different 3D coordinates. So, there are typically as many points in the point cloud from the same laser beam as there are returns. And the **return number** (**return_num**) gives information, which return it was.

First, we take a look at the data types.

```
[9]: print('Intensity:', file.intensity.dtype)
    print('Number of returns:', file.num_returns.dtype)
    print('Return number:', file.return_num.dtype)
```

Intensity: uint16

Number of returns: uint8 Return number: uint8

Since the data types are integer values, we can also take a look at how many unique values there are.

```
[10]: print('Intensity:', np.unique(file.intensity))
print('Number of returns:', np.unique(file.num_returns))
print('Return number:', np.unique(file.return_num))
```

```
Intensity: [
                        2
                            3
                                         6
                                             7
                                                  8
                                                         10
                                                                  12
                   1
                                                      9
                                                              11
                                                                       13
                                                                           14
                                                                               15
                                                                                    16
17
                                                                   33
  18
      19
               21
                   22
                        23
                            24
                                25
                                     26
                                         27
                                             28
                                                  29
                                                      30
                                                          31
                                                               32
                                                                        34
                                                                            35
          20
  36
      37
           38
               39
                   40
                        41
                            42
                                43
                                     44
                                         45
                                             46
                                                  47
                                                      48
                                                           49
                                                               50
                                                                   51
                                                                        52
                                                                            53
      55
                                                                            71
  54
          56
               57
                   58
                        59
                            60
                                61
                                     62
                                         63
                                             64
                                                  65
                                                      66
                                                           67
                                                               68
                                                                   69
                                                                        70
  72
      73
          74
               75
                   76
                        77
                                     80
                                         81
                                             82
                                                  83
                                                      84
                                                           85
                                                               86
                                                                   87
                            78
                                79
                                                                            89
  90
      91
          92
               93
                   94
                        95
                            96
                                97
                                     98
                                         99 100 101 102 103 104 105 106 107
 108 109 110 111 112 113 114 115 116 117 118 119 120 121 122 123 124 125
 126 127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143
 144 145 146 147 148 149 150 151 152 153 154 155 156 157 158 159 160 161
 162 163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179
 180 181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197
 198 199 200 201 202 203 204 205 206 207 208 209 210 211 212 213 214 215
 216 217 218 219 220 221 222 223 224 225 226 227 228 229 230 231 232 233
```

```
234 235 236 237 238 239 240 241 242 243 244 245 246 247 248 249 250 251 252 253 254 255]
```

Number of returns: [1 2 3 4] Return number: [1 2 3 4]

2.2 Intensity

Because of the large number of different values for intensity, we use a bar plot to show how many times these values occur. (We could also use a histogram plot, but then would need to define the bins ourselves.) The parameter **width** determines the occupied bar width, and a value of 1.0 means that the bars are one after another without any gaps. With this many bars, not all gaps would be visible and it would make the visualization look unattractive.

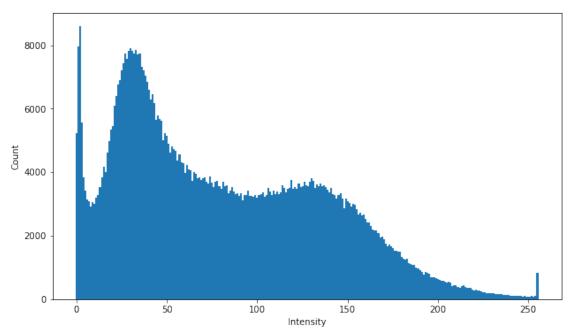
```
[11]: u, c = np.unique(file.intensity, return_counts=True)

fig = plt.figure(figsize=(8, 5))
  fig.suptitle('Intensity Values', fontsize=14, fontweight='semibold')

ax = fig.add_axes((0, 0, 1, 0.90))
  ax.bar(u, c, width=1.0)
  ax.set_xlabel('Intensity')
  ax.set_ylabel('Count')
```

[11]: Text(0, 0.5, 'Count')

Intensity Values

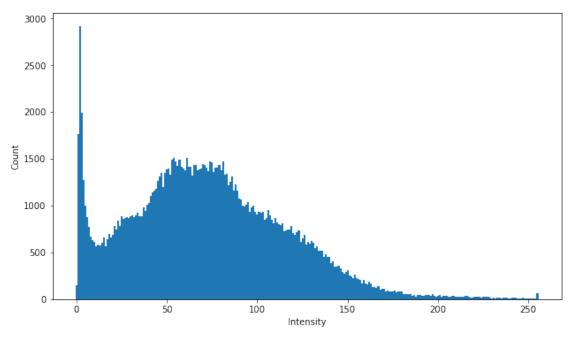


Let us show the intensity only for one class. With the NumPy function **where()**, we can get an array of indices, where a given condition is true. By subsequently applying the function **take()** with these indices, the values can be taken at the appropriate indexes.

What follows are the intensities for the class roof.

[12]: Text(0, 0.5, 'Count')

Intensity Values of Roof Points



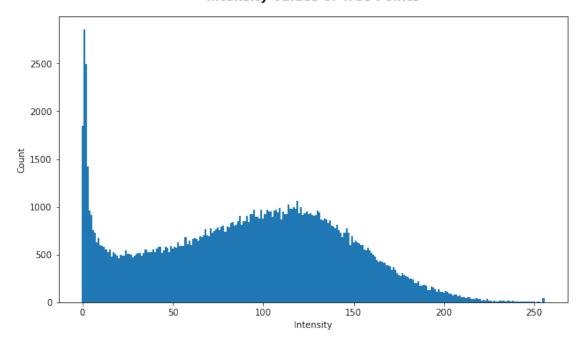
And here for the class trees.

```
[13]: tree_intensities = np.take(file.intensity, np.where(classes == 8))

u, c = np.unique(tree_intensities, return_counts=True)
```

[13]: Text(0, 0.5, 'Count')

Intensity Values of Tree Points



We can already see that the roof points have generally a lower intensity as the tree points.

Next, we plot several bar plots on top of each other, one per class, so that we can better see which object categories emit a low or high intensity return of the laser pulse. For this purpose, we loop over all class names, extract all points of a certain class, construct a bar plot on the axes, and give the plot a label (with the **label** parameter) that is used create a legend for the axes with the method **legend()**. We let matplotlib find the best location for the legend.

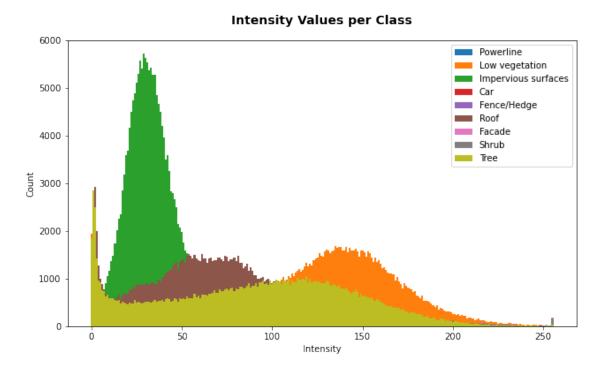
```
[14]: fig = plt.figure(figsize=(8, 5))
  fig.suptitle('Intensity Values per Class', fontsize=14, fontweight='semibold')
  ax = fig.add_axes((0, 0, 1, 0.90))
```

```
ax.set_xlabel('Intensity')
ax.set_ylabel('Count')

for i,n in enumerate(class_names):
    u, c = np.unique(np.take(file.intensity, np.where(classes == i)),
    return_counts=True)
    ax.bar(u, c, width=1.0, label=n)

ax.legend(loc='best')
```

[14]: <matplotlib.legend.Legend at 0x7f443b9c1b90>



Although the underrepresented classes are not really visible, a general trend can be seen for the 4 classes that are well represented. Impervious surfaces mostly emit low intensity, roof surfaces higher intensity, trees seem to be more balanced with a peak between roof surfaces and low vegetation, and low vegetation has mostly higher intensity.

If we wanted to assign points to categories (perform a classification by hand) based only on their intensity, then we could try to find values for intensity that separates the individual class ranges. For example, a point with intensity <= 50 is more likely to belong to the impervious surface class than to the low vegetation class (or the other 2 mentioned classes). A machine learning algorithm would find these separating values in such a way that the correct class is predicted for most points. If we then have a point without a class, we could check in which range this point falls according to its intensity, and give it the respective class label. However, the task of classification is rarely based on one feature only, but more often on several features.

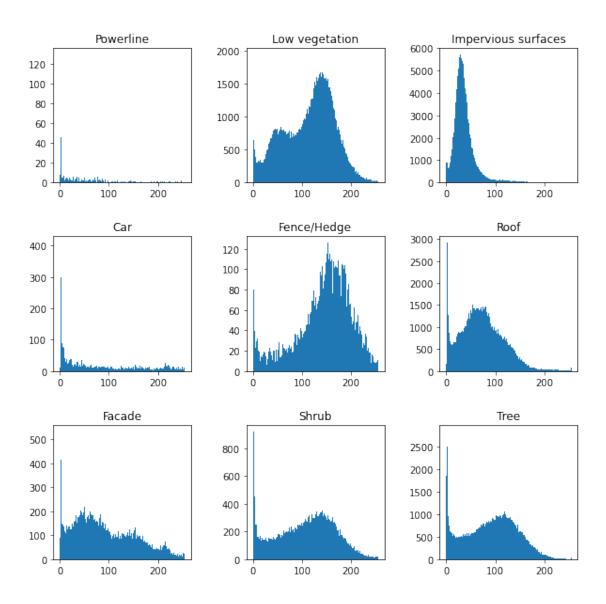
The next example shows subplots, in which all the classes are given in their own plot. The function **subplots()** constructs a figure and a number of axis according to the given number of rows (first parameter) and number of columns (second parameter). We can also define that the properties (scale, ticks, etc.) of the axes objects can be shared for the x-axis (**sharex**) and y-axis (**sharey**) in each column or row, respectively. Try also the value **all**, so that the scales are the same for all plots.

Because the plots would be rather close together, we adjust the space between subplots with the **subplots_adjust()** method. The two parameters defined are given as a fraction of the average axis width or height reserved for the space between subplots. (Try to comment out this line to see how it would otherwise look like.)

The axes object (returned by subplots) is a 2D array of the number of axes constructed, which need to be indexed accordingly with 2 index values. Here, we calculate the row index of the array by dividing the class index (which is the variable i in the loop) by 3 (, because we have 3 plots per row). In order to have an integer number and not a floating point number as the result of the division, we use the floor division operator // instead of the true division (/). (A floating point number as an index value would be considered an error.) For the column index (the second index), we use the modulo operator %, again by 3 for the same reason as before (number of columns per row is 3).

The title for each axes can be set with the **set()** method, where we have to specify which properties of the bar plot we would like to change.

Intensity



What can be learned from the bar plots?

- For all classes, there are quite a number of points without (or zero) intensity values.
- All vegetation related classes like tree, shrub, hedge, and low vegetation have a peak somewhere between 120 to 170. In comparison, all classes that are constructed from building material (stone, concrete, asphalt) like impervious, facade, and roof surfaces, show lower intensity. This means they absorb more light then the vegetation classes.
- Besides the main peak, low vegetation shows another, lower peak at around 50, which is close to the main peak of impervious surfaces. So, it could well be that low vegetation also contains some points from impervious surfaces. Or that the reflections are a mixture of vegetation and ground. This is quite understandable, because signals resulting from these two classes

are difficult to separate. And it seems that rather the vegetation class is "impure" and the impervious surface class is not.

2.3 Number of returns

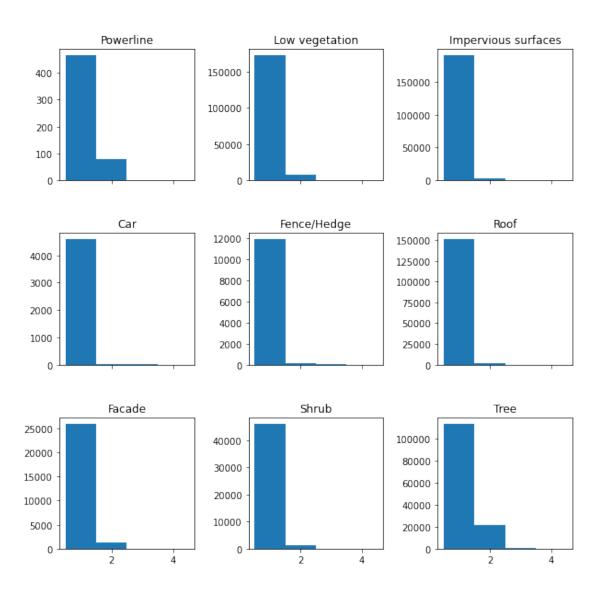
Let us move on to the number of returns. First, we count the occurrences of the 4 values.

```
[16]: u, c = np.unique(file.num_returns, return_counts=True)
print(u, c)
```

```
[1 2 3 4] [716755 36216 902 3]
```

Most laser pulses resulted in only 1 return. Let us check which classes are responsible for it. Maybe some underrepresented class is the only cause.

Number of Returns



Although mostly power lines, trees, and maybe facades resulted in several returns, the higher number of returns are not dominant for these classes. The feature is therefore very likely not a good one to separate any classes.

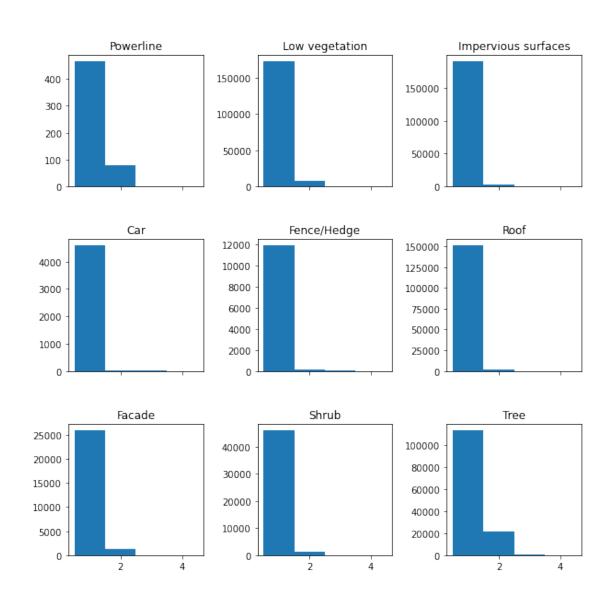
2.4 Return number

As a last attribute of the LAS file, we count the occurrences of the return number.

```
[1 2 3 4] [735339 18147 389 1]
```

And have the same plot as above.

Return Number



From the point properties of the point cloud, only the intensity shows a relevant spectrum of values that can also be allocated to the individual classes. Although with intensity alone, it will only be possible to differentiate between the broad categories of vegetation and construction materials.

We therefore continue with the extracted features.

3 Features

Load the extracted features from the exercise notebook '02 - Feature Calculations with NumPy (Part 2)'.

The following code cell uses the reference features from the '02 - Check Your Features" notebook. If you want to use your own calculated features, then copy them into the same folder as this notebook and uncomment the respective line.

```
[20]: data_dir = str(Path.home()) + r'/coursematerial/GIS/ISPRS/PointsWithFeatures'
filepath = os.path.join(data_dir, r'FeaturesReference.npy')
features = np.load(filepath)
#features = np.load('Features.npy')
print(features.shape)
```

(753876, 23)

Construct a pandas DataFrame, so that we can select the columns with a string index for convenience. Otherwise we would need to remember which column index a certain features has. When selecting a column of a DataFrame, we get an object of type Series, which we can convert with the method **to_numpy()** to a 1D NumPy array with n elements to use in matplotlib. However, matplotlib also accepts pandas DataFrames as data.

```
'std_z',
                        'sum_of_eigenvalues_2d',
                        'ratio_of_eigenvalues_2d',
                        'radius_knn_2d',
                        'density_2d',
                        'eigenvalue1',
                        'eigenvalue2',
                        'eigenvalue3',
                        'eigenvalue2D1',
                        'eigenvalue2D2'
                       ٦
      df = pd.DataFrame(features, columns=feature names)
      df
[21]:
                          planarity
                                     scattering
              linearity
                                                  omnivariance
                                                                 anisotropy \
      0
               0.514636
                           0.484771
                                        0.000593
                                                      0.044426
                                                                   0.999407
      1
               0.286950
                           0.712251
                                        0.000799
                                                      0.048378
                                                                   0.999201
      2
               0.337194
                           0.662073
                                        0.000733
                                                      0.047258
                                                                   0.999267
      3
               0.670130
                           0.327899
                                        0.001971
                                                      0.065043
                                                                   0.998029
               0.591237
                           0.406962
                                        0.001801
                                                      0.064009
                                                                   0.998199
                                       0.003587
                                                      0.077360
                                                                   0.996413
      753871
               0.738930
                           0.257483
                                       0.001024
                                                      0.053272
      753872
               0.466459
                           0.532517
                                                                   0.998976
      753873
               0.730894
                           0.266702
                                        0.002404
                                                      0.068023
                                                                   0.997596
      753874
               0.718838
                           0.277748
                                        0.003415
                                                      0.076797
                                                                   0.996585
      753875
               0.781174
                           0.215357
                                        0.003469
                                                      0.074632
                                                                   0.996531
              eigenentropy
                             sum_eigenvalues
                                               change_of_curvature radius_knn3d \
      0
                   0.635132
                                    0.514617
                                                          0.000399
                                                                         1.423552
      1
                  0.682779
                                    0.483546
                                                          0.000466
                                                                         1.291007
      2
                   0.675992
                                    0.485428
                                                          0.000441
                                                                         1.257975
      3
                   0.570471
                                    0.566638
                                                          0.001480
                                                                         1.278241
      4
                   0.611306
                                    0.947597
                                                          0.001277
                                                                         1.804384
                                                          0.002836
                                                                         0.511273
      753871
                   0.528006
                                    0.093828
      753872
                   0.651261
                                    0.097857
                                                          0.000667
                                                                         0.512640
      753873
                   0.529421
                                    0.097016
                                                          0.001891
                                                                         0.487750
      753874
                   0.543241
                                    0.093028
                                                          0.002658
                                                                         0.521632
      753875
                   0.488836
                                    0.111733
                                                          0.002838
                                                                         0.520769
              density_3d
                                 std_z sum_of_eigenvalues_2d \
      0
                1.655089
                              0.035433
                                                      0.513361
      1
                2.218989
                              0.035964
                                                      0.482252
      2
                2.398423 ...
                              0.035526
                                                      0.484166
```

0.563316

3

2.286142 ...

0.057638

4	0.812746	0.070328	0	.942651		
•••	••• •••	•••	•••			
753871	35.725886	0.022804		.093308		
753872	35.440786	0.016182		.097595		
753873	41.148094	0.018638	0	.096668		
753874	33.639426	0.019628	0	.092643		
753875	33.807018	0.023458	0.111183			
	ratio_of_eig	genvalues_2d	radius_knn_2d	density_2d	eigenvalue1	\
0		0.483495	1.423517	3.141629	0.672967	
1		0.711442	1.290039	3.825380	0.583482	
2		0.663891	1.257816	4.023891	0.601128	
3		0.331784	1.273499	3.925390	0.750840	
4		0.410825	1.799694	1.965543	0.708937	
		•••	•••	•••	•••	
753871		0.260700	0.510392	24.438379	0.790728	
753872		0.531483	0.511077	24.372886	0.651650	
753873		0.268709	0.487647	26.771227	0.786466	
753874		0.281240	0.520096	23.534927	0.778466	
753875		0.218782	0.519230	23.613493	0.818133	
	eigenvalue2	eigenvalue3	eigenvalue2D1	eigenvalue	2D2	
0	0.326634	0.000399	0.674084	0.325	916	
1	0.416052	0.000466	0.584302	0.415	698	
2	0.398431	0.000441	0.601001	0.398	999	
3	0.247680	0.001480	0.750873	0.249	127	
4	0.289787	0.001277	0.708805	0.291	195	
•••	•••	•••	•••	•••		
753871	0.206435	0.002836	0.793210	0.206	790	
753872	0.347682	0.000667	0.652962	0.347	038	
753873	0.211643	0.001891	0.788203	0.211	797	
753874	0.218875	0.002658	0.780494	0.219	506	
753875	0.179029	0.002838	0.820491	0.179	509	

[753876 rows x 23 columns]

Add the classification of the points to the features DataFrame, so that everything is together.

```
[22]: df['class'] = classes
df
```

[22]:	linearity	planarity	scattering	omnivariance	anisotropy	\
0	0.514636	0.484771	0.000593	0.044426	0.999407	
1	0.286950	0.712251	0.000799	0.048378	0.999201	
2	0.337194	0.662073	0.000733	0.047258	0.999267	
3	0.670130	0.327899	0.001971	0.065043	0.998029	

4	0.591237	0.406962	0.0018	01	0.064009	0.99819	99	
•••	•••		•••		•••			
753871	0.738930	0.257483	0.0035	87	0.077360	0.99643	13	
753872	0.466459	0.532517	0.0010		0.053272	0.99897	76	
753873	0.730894	0.266702	0.0024	.04	0.068023	0.99759	96	
753874	0.718838	0.277748	0.0034	15	0.076797	0.99658	35	
753875	0.781174	0.215357	0.0034	69	0.074632	0.99653	31	
			_	_	_			
	eigenentrop	•	genvalues	change_	of_curvature		ıs_knn3d	\
0	0.63513		0.514617		0.000399		1.423552	
1	0.68277		0.483546		0.000466		1.291007	
2	0.67599		0.485428		0.00044		1.257975	
3	0.57047		0.566638		0.001480		1.278241	
4	0.61130	06	0.947597		0.00127	7 :	1.804384	
•••	•••		•••		•••	•••		
753871	0.52800		0.093828		0.00283		0.511273	
753872	0.65126		0.097857		0.00066		0.512640	
753873	0.52942	21	0.097016		0.00189		0.487750	
753874	0.54324	1	0.093028		0.002658	3 (0.521632	
753875	0.48883	86	0.111733		0.002838	3 (0.520769	
	density_3d	sum_of	_eigenval		ratio_of_ei	-		
0	1.655089	•••		513361			33495	
1	2.218989	•••	0.	482252		0.73	11442	
2	2.398423	•••	0.	484166		0.66	53891	
3	2.286142	•••	0.	563316		0.33	31784	
4	0.812746		0.	942651		0.43	10825	
•••			•••	•		•••		
753871	35.725886	•••	0.	093308		0.26	30700	
753872	35.440786	•••	0.	097595		0.53	31483	
753873	41.148094	•••	0.	096668		0.26	58709	
753874	33.639426	•••	0.	092643		0.28	31240	
753875	33.807018		0.	111183		0.23	18782	
	1: 1	0.1 . 1	01			- 0		,
^	radius_knn_ 1.4235			envalue1	•	•	envalue3	\
0				0.672967			0.000399	
1	1.2900			0.583482			0.000466	
2	1.2578			0.601128			0.000441	
3	1.2734			0.750840			0.001480	
4	1.7996	1.96	35543	0.708937	0.28978	37 (0.001277	
 753871	0.5103	 892 24.43	 18370	0.790728	0.20643	 35 (0.002836	
753871	0.5103			0.790728			0.002636	
753873	0.3110			0.786466			0.000887	
753874	0.5200			0.778466			0.002658	
753875	0.5192	23.61	.3493	0.818133	0.1790	29 (0.002838	

	eigenvalue2D1	eigenvalue2D2	class
0	0.674084	0.325916	1
1	0.584302	0.415698	1
2	0.601001	0.398999	1
3	0.750873	0.249127	1
4	0.708805	0.291195	1
	•••		
753871	0.793210	0.206790	2
753872	0.652962	0.347038	2
753873	0.788203	0.211797	2
753874	0.780494	0.219506	2
753875	0.820491	0.179509	2

[753876 rows x 24 columns]

The method **describe()** of Pandas shows some basic statistics that can be useful for a first look at the data.

(See how the mean of classes makes no real sense, but the 50% percentile (or median) is more meaningful.)

[23]: df.describe()

[23]:		linearity	planarity	scattering	omnivariance	\	
	count	753876.000000	753876.000000	753876.000000	753876.000000	•	
	mean	0.557764	0.387218	0.055018	0.124257		
	std	0.254856	0.251118	0.093513	0.088348		
	min	0.000276	0.000062	0.000008	0.003198		
	25%	0.349201	0.169167	0.001215	0.051112		
	50%	0.555117	0.363175	0.005859	0.079920		
	75%	0.774457	0.590015	0.074130	0.204756		
	max	0.999620	0.999046	0.869420	0.332718		
		anisotropy	eigenentropy	sum_eigenvalues	change_of_cu	rvature	\
	count	753876.000000	753876.000000	753876.000000	753876	.000000	
	mean	0.944982	0.637717	0.530745	5 0	.032339	
	std	0.093513	0.216737	0.644467	0	.049539	
	min	0.130580	0.004297	0.000411	. 0	.000007	
	25%	0.925870	0.537542	0.179740	0	.000844	
	50%	0.994141	0.673145	0.362666	0	.004596	
	75%	0.998785	0.768934	0.676156	0	.049049	
	max	0.999992	1.096792	34.571103	0	.307947	
		radius_knn3d	density_3d	sum_of_eiger	values_2d \		
	count	753876.000000	753876.000000	0000 753876.000000			
	mean	0.972458	118.045849		0.434099		
	std	0.487351	1200.733030	•••	0.441575		
	min	0.030000	0.002944	•••	0.000048		

```
25%
             0.671118
                             2.419261
                                                         0.157143
50%
             0.890730
                             6.756209
                                                         0.335010
75%
             1.254352
                            15.795880
                                                         0.515722
            11.749357
                       176838.830873
                                                        27.902223
max
                                                     density_2d
       ratio_of_eigenvalues_2d
                                                                     eigenvalue1
                                 radius_knn_2d
                  753876.000000
                                  753876.000000
                                                  753876.000000
                                                                  753876.000000
count
                       0.411878
                                        0.935126
                                                      44.341698
                                                                        0.694408
mean
                       0.262245
                                                     467.636812
                                        0.458810
                                                                        0.139602
std
min
                       0.000014
                                        0.014142
                                                        0.067502
                                                                        0.351351
25%
                       0.177836
                                        0.650000
                                                        4.369688
                                                                        0.582220
50%
                       0.406867
                                        0.870919
                                                       8.393141
                                                                        0.664603
75%
                       0.628703
                                        1.207021
                                                      15.067924
                                                                        0.797438
max
                       0.999513
                                       9.711364
                                                   31830.989300
                                                                        0.999530
         eigenvalue2
                          eigenvalue3
                                        eigenvalue2D1
                                                        eigenvalue2D2
       753876.000000
                       753876.000000
                                       753876.000000
                                                        753876.000000
count
             0.273253
                                             0.733804
                                                             0.266196
mean
                             0.032339
std
             0.127630
                             0.049539
                                             0.139479
                                                             0.139479
             0.000380
                             0.000007
                                             0.500122
                                                             0.000014
min
25%
             0.178383
                             0.000844
                                             0.613985
                                                             0.150985
50%
                                                             0.289201
             0.291794
                             0.004596
                                             0.710799
75%
             0.378779
                             0.049049
                                             0.849015
                                                             0.386015
                             0.307947
                                                             0.499878
max
             0.499762
                                             0.999986
                class
count
       753876.000000
             3.937998
mean
             2.642773
std
             0.000000
min
25%
             2.000000
50%
             3.000000
75%
             6.000000
max
             8.000000
```

[8 rows x 24 columns]

3.1 Height

As a first feature, we take a look at the elevation (absolute height, z-coordinate) stored with the points. We can again use the **describe()** method of pandas, which is now applied to a column of type Series.

```
[24]: df['absolute_height'].describe()
```

```
[24]: count
                753876.000000
      mean
                   268.098009
      std
                     5.963596
                   250.600000
      min
      25%
                   264.300000
      50%
                   267.830000
      75%
                   271.840000
      max
                   289.910000
```

Name: absolute_height, dtype: float64

Assuming that there might be outliers to both sides, let us define maybe more interesting percentiles with the **describe()** method. In the following example, we want to additionally output the 1 and 99 percentile to see how much differences the lowest and highest 1% of the points make.

```
[25]: df['absolute_height'].describe(percentiles=[0.01, 0.25, 0.5, 0.75, 0.99])
```

```
[25]: count
                753876.000000
                   268.098009
      mean
      std
                     5.963596
      min
                   250.600000
      1%
                   254.240000
      25%
                   264.300000
      50%
                   267.830000
      75%
                   271.840000
      99%
                   282.250000
                   289.910000
      max
```

Name: absolute_height, dtype: float64

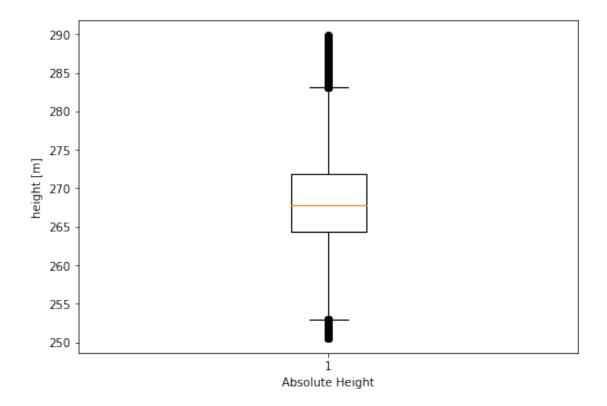
We can either convert the pandas DataFrame into a NumPy array with the method **to_numpy()** to use with matplotlib.

```
[26]: heights = df['absolute_height'].to_numpy()

fig = plt.figure()

ax = fig.add_axes([0, 0, 1, 1])
ax.boxplot(heights)
ax.set_xlabel('Absolute Height')
ax.set_ylabel('height [m]')
```

[26]: Text(0, 0.5, 'height [m]')

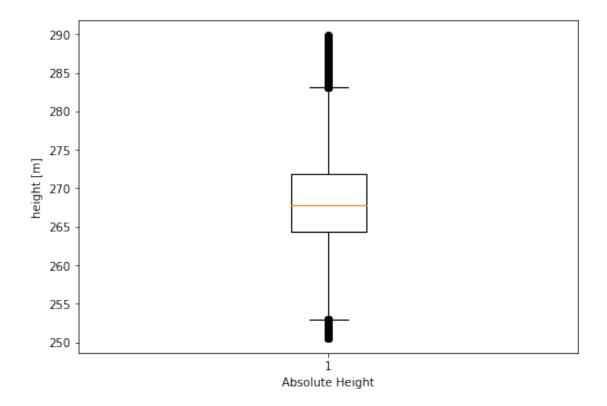


Or use directly a column of the pandas DataFrame.

```
[27]: fig = plt.figure()

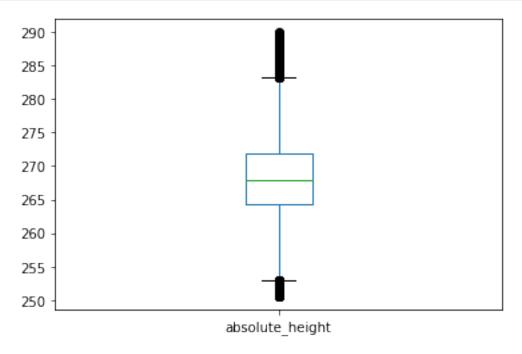
ax = fig.add_axes([0, 0, 1, 1])
ax.boxplot(df['absolute_height'])
ax.set_xlabel('Absolute Height')
ax.set_ylabel('height [m]')
```

[27]: Text(0, 0.5, 'height [m]')



Pandas also offers convenient methods that can be directly called on a column of a DataFrame (which is an object of class Series).





Next, we plot a combination of a box plot and a histogram of heights.

The bin widths are provided with the **bins** parameter as a NumPy array. Here, an array is constructed with the function **arange()** that contains values between start (first parameter), end (second parameter), and with a step size (third parameter) of 0.5. The start and end of the range is the floor of the minimum value and the ceiling of the maximum value in the elevation (height) column.

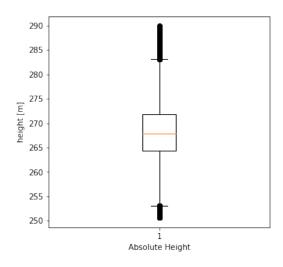
```
[29]: np.arange(np.floor(np.min(heights)), np.ceil(np.max(heights)), 0.5)

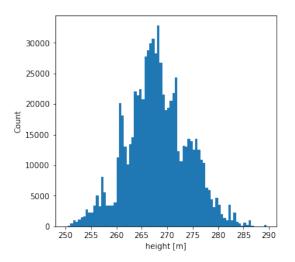
[29]: array([250., 250.5, 251., 251.5, 252., 252.5, 253., 253.5, 254., 254.5, 255., 255.5, 256., 256.5, 257., 257.5, 258., 258.5, 259., 259.5, 260., 260.5, 261., 261.5, 262., 262.5, 263., 263.5, 264., 264.5, 265., 265.5, 266., 266.5, 267., 267.5, 268., 268.5, 269., 269.5, 270., 270.5, 271., 271.5, 272., 272.5, 273., 273.5, 274., 274.5, 275., 275.5, 276., 276.5, 277., 277.5, 278., 278.5, 279., 279.5, 280., 280.5, 281., 281.5, 282., 282.5, 283., 283.5, 284., 284.5, 285., 285.5, 286., 286.5, 287., 287.5, 288., 288.5, 289., 289.5])
```

These bins are then provided with the parameter **bins** of the **hist()** method to construct a histogram.

[30]: Text(0, 0.5, 'Count')

Elevation





And finally, a histogram of heights per classes as shown above for the intensity values.

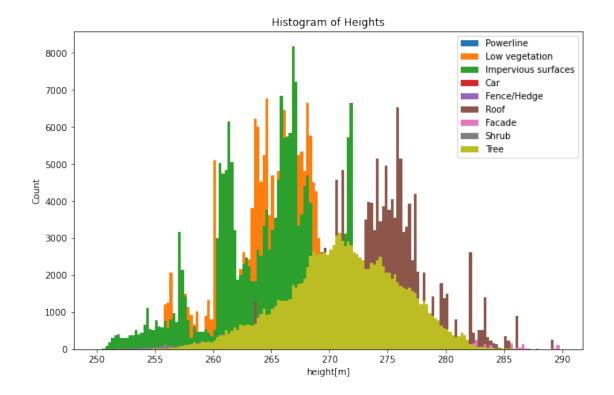
```
fig = plt.figure(figsize=(8, 5))
ax = fig.add_axes((0, 0, 1, 1))

bins = np.arange(np.floor(np.min(heights)), np.ceil(np.max(heights)), 0.25)

for i,n in enumerate(class_names):
    ax.hist(heights[np.where(classes == i)], bins=bins)

ax.legend(labels=class_names)
ax.set_title('Histogram of Heights')
ax.set_xlabel('height[m]')
ax.set_ylabel('Count')
```

[31]: Text(0, 0.5, 'Count')

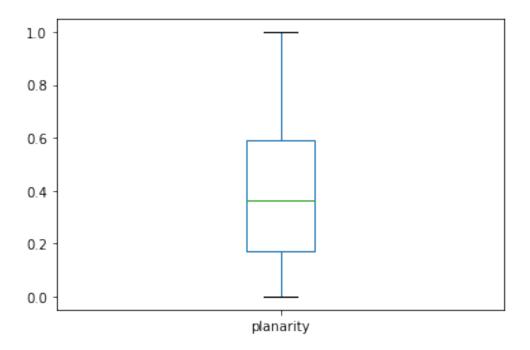


Take a look at these heights and their classes and think if they make sense or not. It might also help to take a look at the point cloud in a 3D visualization, e.g., using CloudCompare.

3.2 Planarity

Next, we turn to the feature of planarity. One could assume that impervious ground and roof surfaces have a high planarity and that vegetation has a rather low planarity.

We start with a simple boxplot of planarity values, and for simplicity use the pandas method **box()**.



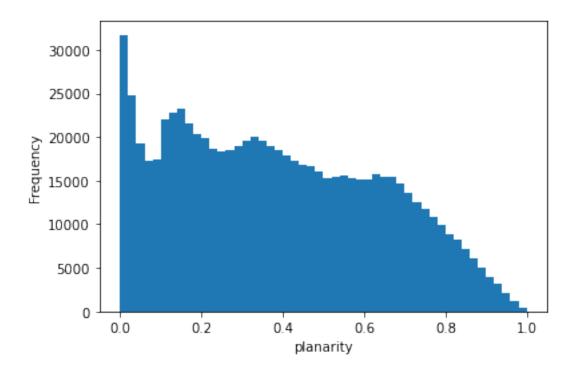
Pandas also has a quick way to get the percentiles or quantiles with the **percentile()** or **quantile()** method, respectively.

We can now see where most of the values are located. The median is not quite at 0.5, but also not that far off. There are hardly any perfect planarity values.

Let us also plot a histogram of planarity values. Note that the pandas method returns an Axes object, which we can use to set some properties of the plot.

```
[34]: ax = df['planarity'].plot.hist(bins=50)
ax.set_xlabel('planarity')
```

[34]: Text(0.5, 0, 'planarity')



Next, a histogram plot per class is generated, kind of like we did with the bar plots.

```
[35]: planarity = df['planarity'].to_numpy()

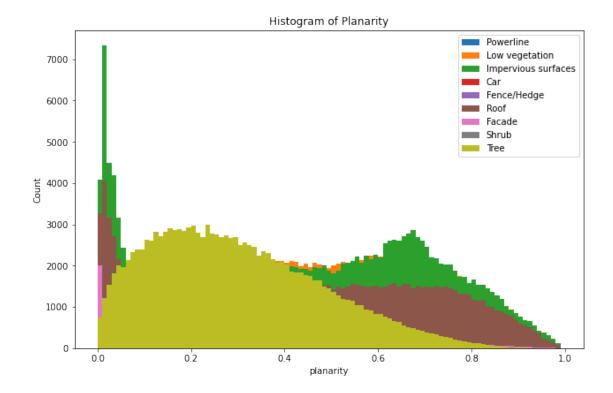
fig = plt.figure(figsize=(8, 5))
ax = fig.add_axes((0, 0, 1, 1))

bins = np.arange(0.0, 1.0, 0.01)

for i,n in enumerate(class_names):
    ax.hist(planarity[np.where(classes == i)], bins=bins)

ax.legend(labels=class_names)
ax.set_title('Histogram of Planarity')
ax.set_xlabel('planarity')
ax.set_ylabel('Count')
```

[35]: Text(0, 0.5, 'Count')



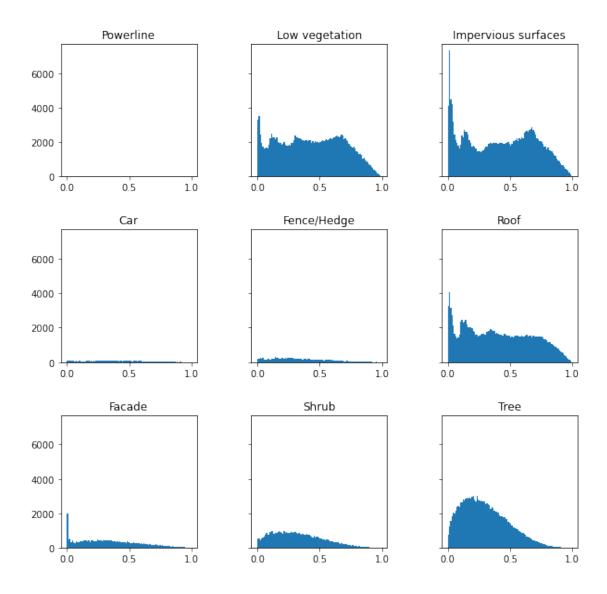
Unfortunately, the tree class occludes almost everything. But it can be seen that the tree class has mostly low planarity.

So let us better plot all the classes on their own.

```
fig, ax = plt.subplots(3, 3, sharex='none', sharey='all', figsize=(10, 10))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
fig.suptitle('Histograms of Planarity', fontsize=14, fontweight='bold')

for i,n in enumerate(class_names):
    ax[i//3, i%3].hist(planarity[np.where(classes == i)], bins=np.arange(0.0, 1.
    -0, 0.01))
    ax[i//3, i%3].set(title=class_names[i])
```

Histograms of Planarity



As a conclusion, the tree, shrub, fence/hedge classes clearly have a tendency to lower planarity. And although impervious surfaces and roofs have more points with high planarity values, they are nevertheless rather homogeneously distributed. Perfectly planar points are very seldom, which might have to do with the low density of the point cloud itself, and the comparably large neighborhood of 20 points from which the feature is calculated. All 20 points would need to be located on a plane for a very high planarity value.

The point cloud has a point density of approximately 4 points per square meter. And a group of 20 points therefore take over 4 square meters. There are not that many points with that many neighbor points that make flat areas.

By the way, to better see the histograms of classes like power line, car, and fence/hedge, you can

set the **sharey** parameter to none in the above figure.

4 Scattering

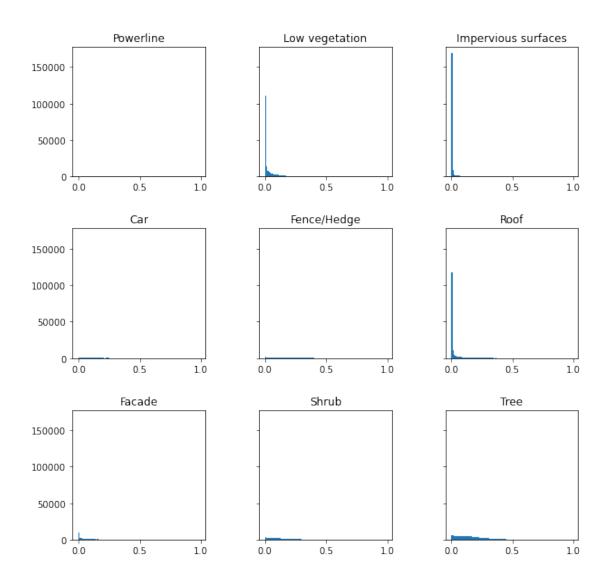
For the feature, scattering, we would expect a reversed image. So, let us start with the histograms per classes.

```
[37]: scattering = df['scattering'].to_numpy()

fig, ax = plt.subplots(3, 3, sharex='none', sharey='all', figsize=(10, 10))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
fig.suptitle('Histograms of Scattering', fontsize=14, fontweight='bold')

for i,n in enumerate(class_names):
    ax[i//3, i%3].hist(scattering[np.where(classes == i)], bins=np.arange(0.0, 0.0))
ax[i//3, i%3].set(title=class_names[i])
```

Histograms of Scattering

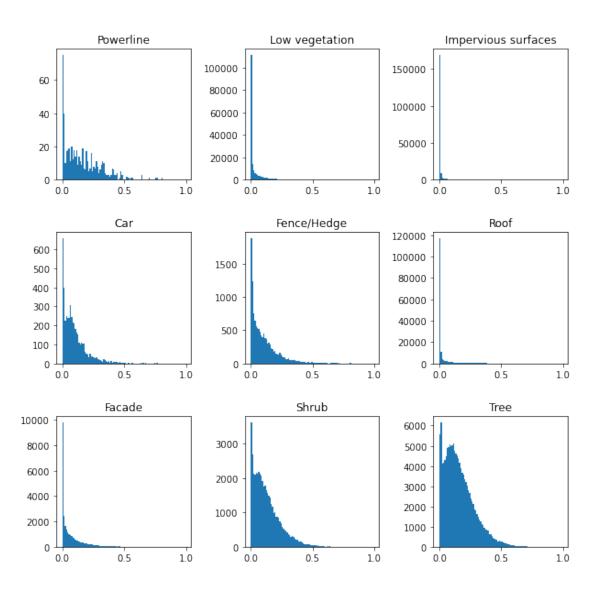


That is an unexpected outcome. It can be improved by turning off to option to share the y-scale.

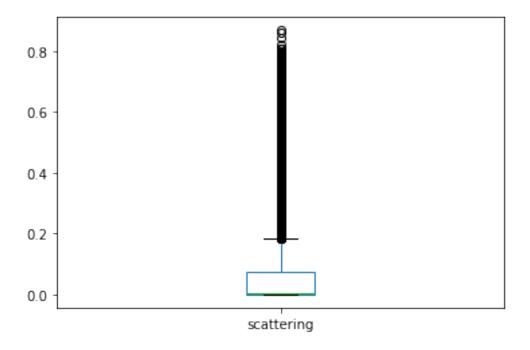
```
fig, ax = plt.subplots(3, 3, sharex='none', sharey='none', figsize=(10, 10))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
fig.suptitle('Histograms of Scattering', fontsize=14, fontweight='bold')

for i,n in enumerate(class_names):
    ax[i//3, i%3].hist(scattering[np.where(classes == i)], bins=np.arange(0.0, u)
    -1.0, 0.01))
    ax[i//3, i%3].set(title=class_names[i])
```

Histograms of Scattering



To investigate this further, let us take a look at a box plot.



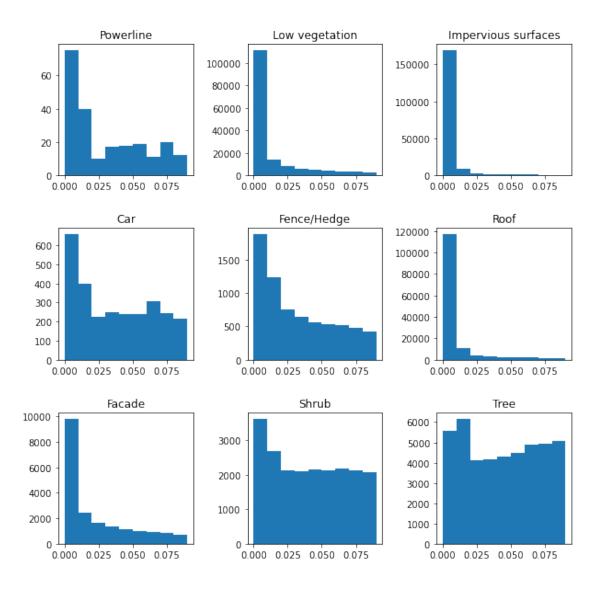
It seems that most points have a low scattering value. But at least from the second histogram plot, we can see that tree, shrub, and fence/hedge do actually have some noticeable scattering that is higher than 0.025. We can concentrate on the values with low scattering (<0.01) and have another plot.

Notice how the two conditions are combined with the $logical_and()$ function of NumPy. Each condition ("classes == i" and "scattering < 0.1") gives an array of Boolean values, that have the value true where the condition is true. These two arrays of Boolean values are then combined to a single Boolean array, which can ten be used in where() to find the indices where both conditions hold.

```
[40]: fig, ax = plt.subplots(3, 3, sharex='none', sharey='none', figsize=(10, 10))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
fig.suptitle('Histograms of Scattering', fontsize=14, fontweight='bold')

for i,n in enumerate(class_names):
    ax[i//3, i%3].hist(scattering[np.where(np.logical_and(classes == i,u)))
    scattering < 0.1))], bins=np.arange(0.0, 0.1, 0.01))
    ax[i//3, i%3].set(title=class_names[i])</pre>
```

Histograms of Scattering



From this we notice that impervious surfaces and roof have a very low scattering as expected, because they should mostly consists of flat surfaces. (Although a more pronounced value range would have been nicer.) Because only few beams of an aerial laser scanner hits building facades, these points are also rather scattered instead of planar. This is to be expected.

4.1 Density 3D

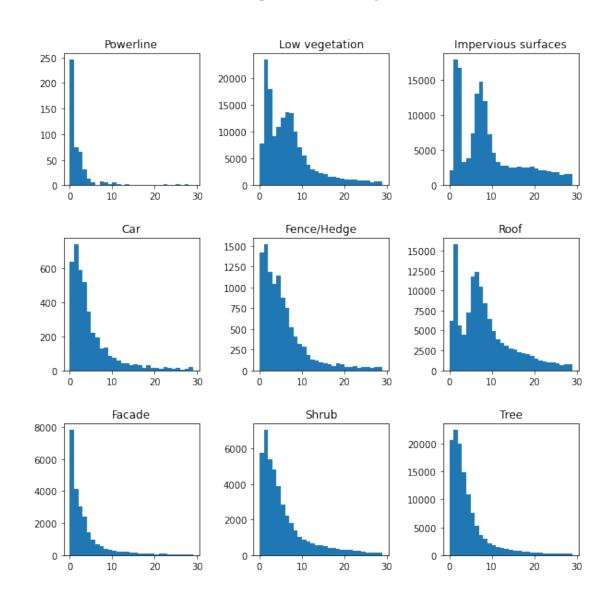
As a last feature, we take a look at the 3d density of the k nearest neighbor points. We limit the range of bins to be between 0 and 30 as most values are in this range.

```
[41]: density3d = df['density_3d'].to_numpy()

fig, ax = plt.subplots(3, 3, sharex='none', sharey='none', figsize=(10, 10))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
fig.suptitle('Histograms of Density 3D', fontsize=14, fontweight='bold')

for i,n in enumerate(class_names):
    ax[i//3, i%3].hist(density3d[np.where(classes == i)], bins=np.arange(0.0, u)
    30.0, 1.0))
    ax[i//3, i%3].set(title=class_names[i])
```

Histograms of Density 3D



It is interesting to note that vegetation classes seem to have more values close to 0 and impervious surfaces and roof have higher values.

5 Scatterplots

Before we come to an end, let us also do a few scatterplots in the end. A scatterplot is constructed with the **scatter()** method of an Axes object and takes two NumPy arrays as input. The **s** parameter gives the marker size. As there are many points, a very small marker size is used.

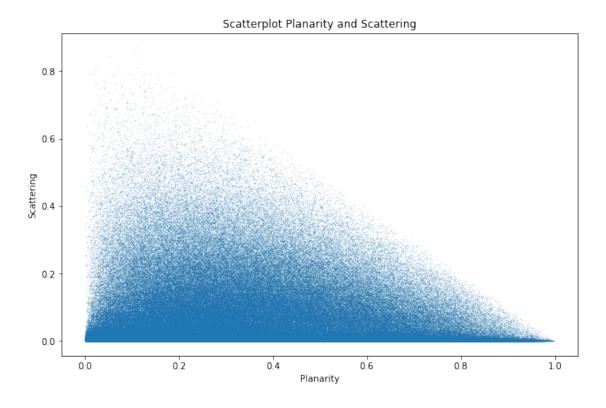
```
[42]: planarity = df['planarity'].to_numpy()
    scattering = df['scattering'].to_numpy()

fig = plt.figure(figsize=(8, 5))
    ax = fig.add_axes((0, 0, 1, 1))

ax.scatter(planarity, scattering, s=0.005)

ax.set_title('Scatterplot Planarity and Scattering')
    ax.set_xlabel('Planarity')
    ax.set_ylabel('Scattering')
```

[42]: Text(0, 0.5, 'Scattering')



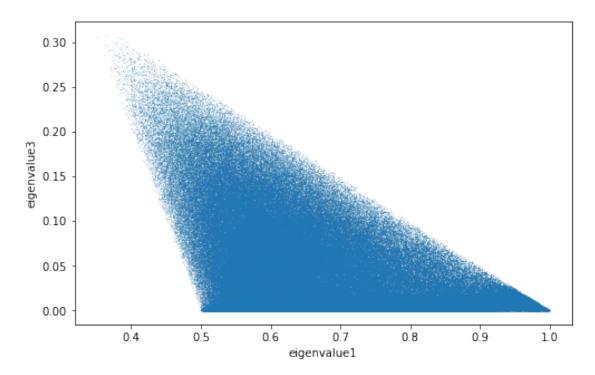
Because most scattering values are rather low, most points gather at the bottom of the plot.

Let us try the eigenvalues 1 and 3.

Pandas has a nice interface where we only need to provide the two column names of the features to plot.

```
[43]: df.plot.scatter(x='eigenvalue1', y='eigenvalue3', figsize=(8, 5), s=0.005)
```

[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4438216650>



As the eigenvalues are dependent on each other, we see some negative correlation, but there is also some differences.

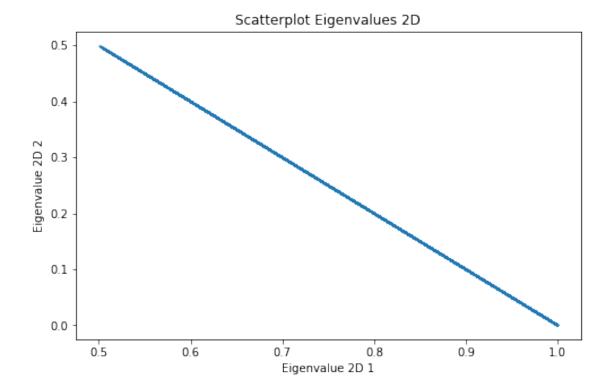
Even more prominent is the correlation of the 2D eigenvalues as they are even defined accordingly. Since the **scatter()** method of pandas returns an Axes object, we can change the plot properties accordingly.

```
[44]: ax = df.plot.scatter(x='eigenvalue2D1', y='eigenvalue2D2', figsize=(8, 5), s=0.

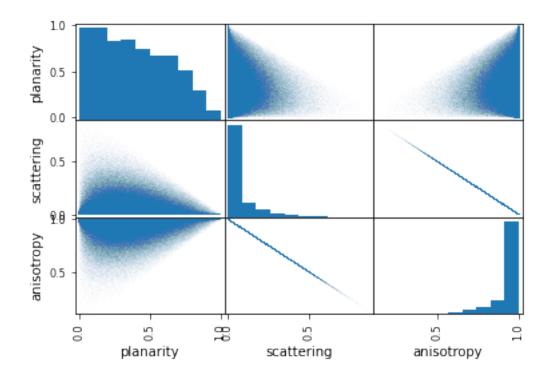
→005)

ax.set_title('Scatterplot Eigenvalues 2D')
ax.set_xlabel('Eigenvalue 2D 1')
ax.set_ylabel('Eigenvalue 2D 2')
```

[44]: Text(0, 0.5, 'Eigenvalue 2D 2')



Pandas provides also a **scatter_matrix()** function, although we need to provide a DataFrame as first argument.



Note that the diagonals are not scatterplots, as they would be scatterplots of the feature with themselves, but rather histograms of the respective features are shown. (As an alternative, kernel density estimations can also be displayed in the diagonal.)

6 Final words

Feel free to explore further features on your own in the same or similar manner from what is shown above.

You can also try a different k in the k nearest neighbor query when calculating the features and see if the features look different.

There are many more ways to customize plots in matplotlib. We leave it to you to further explore them.